

TALARI PRADEEP

AI Infrastructure Engineer | Platform Engineer | DevOps & MLOps Specialist

Bangalore, India • talaripradeep45@gmail.com • +91-7330737488 •

linkedin.com/in/pradeep-talari • github.com/Pradeeptalari14 • talaripradeep.info

PROFESSIONAL SUMMARY

Results-driven **AI Infrastructure & Platform Engineer** with **6+ years of enterprise IT operations experience**, architecting and deploying scalable AI/ML platforms, LLMOps pipelines, and cloud-native DevOps ecosystems. Proven track record delivering enterprise-grade Agentic AI systems, RAG platforms, and multi-agent orchestration frameworks across Azure, AWS, and GCP. Deep expertise in Infrastructure as Code (Terraform, Pulumi, Ansible), container orchestration (Kubernetes, Helm, ArgoCD), and full-stack observability (Prometheus, Grafana, OpenTelemetry). Architected GenAI platforms integrating GPT-4, Claude, Gemini, Llama, and open-source models—driving 40%+ operational efficiency gains and 58% reduction in MTTR. Passionate about building self-healing, AI-augmented infrastructure platforms that accelerate developer velocity and enable enterprise AI transformation at scale.

58% MTTR Reduction via AI monitoring copilot & telemetry correlation

\$300K Annual Savings via VMware compute consolidation & cloud FinOps

99.8% Platform Uptime for multi-agent OS supporting 50+ workflows

94% Lead Time Cut (3 days → 2 hours) via GitOps ArgoCD & Helm

CORE TECHNICAL PROFICIENCIES

AI / LLMS & AGENTS

GPT-4o, Claude 3.5, Gemini 1.5, Llama 3, Mistral, DeepSeek, Ollama, vLLM, LangChain, LangGraph, LlamaIndex, CrewAI, AutoGen, Haystack, DSPy

RAG & VECTOR SEARCH

Vector DBs (Weaviate, Pinecone, Qdrant, Milvus, ChromaDB), OpenSearch, Elasticsearch, Semantic/Hybrid Search, GraphRAG, Knowledge Graphs, Metadata Filtering

AI TECHNIQUES & LLMOps

Prompt Engineering, Function/Tool Calling, Structured Outputs, Embeddings, Chunking, Multi-Agent Systems, MCP (Model Context Protocol), LLMOps, MLOps, AI Governance

CLOUD PLATFORMS

Microsoft Azure, Amazon Web Services (AWS), Google Cloud Platform (GCP) — Multi-Account Landing Zones, IAM Boundaries, Transit Gateway, FinOps Cost Optimization

DEVOPS & IAC

Docker, Kubernetes (AKS, EKS, GKE), Helm, ArgoCD, Terraform, Terragrunt, Pulumi, Ansible, GitHub Actions, Jenkins, GitOps Pipelines

OBSERVABILITY & SRE

Prometheus, Grafana, OpenTelemetry (OTel), Loki, Tempo, LogicMonitor, Alertmanager, Distributed Tracing, SLO/SLI Monitoring & Error Budgets

SYSTEMS & VIRTUALIZATION

Windows Server (2012–2022), VMware vSphere 7.x/8.x (ESXi, vCenter, vSAN), Hyper-V, Dell MX7000, Cisco UCS, NetApp SAN/NAS, PowerShell DSC, Linux (Ubuntu, RHEL)

DATABASES & CACHING

PostgreSQL, pgvector, Redis, MySQL, MongoDB, Neo4j, Vector Databases — Schema Design, High Availability, Query Optimization

LANGUAGES & SCRIPTS

Python (FastAPI, asyncio), Bash, PowerShell, Go, YAML, HCL, SQL

ITSM & GOVERNANCE

ServiceNow ITSM (Incident, Problem, Change, Patch Management, CAB), ITIL v4 processes, SOC2 compliance

VOICE & REALTIME AI

Whisper STT, ElevenLabs TTS, WebRTC, Voice AI Receptionist, Real-Time Audio Streaming, Speech AI

PROFESSIONAL EXPERIENCE

Infrastructure & AI Engineer · Accenture · Bangalore, India

Apr 2022 – Present

- **Enterprise Agentic AI OS:** Architected and deployed an Enterprise Agentic AI Operating System integrating LangGraph multi-agent orchestration, RAG pipelines, and tool-calling capabilities across 50+ enterprise workflows—reducing manual processing time by 65% and achieving 99.8% uptime SLA.
- **Cloud-Native MLOps & GitOps:** Designed and operationalized a cloud-native MLOps/LLMOps platform on Azure AKS using ArgoCD GitOps, Helm charts, and Terraform IaC; automated model promotion pipelines with GitHub Actions, cutting deployment lead time from 3 days to under 2 hours.
- **Enterprise RAG Platform:** Built an Enterprise RAG Platform with Weaviate vector database, OpenSearch hybrid search, and LangChain retrieval chains—delivering semantic document intelligence across 5M+ enterprise knowledge base records with <200ms P99 query latency.
- **AI Incident Copilot:** Engineered an AI Infrastructure Monitoring Copilot using Claude/GPT-4 function calling + LogicMonitor APIs, reducing MTTR by 58% and enabling natural-language incident triage that auto-generated ServiceNow tickets with root cause hypotheses.
- **Multi-Cloud Kubernetes Fleet:** Implemented multi-cloud Kubernetes platform (AKS + EKS + GKE) with Helm umbrella charts, Istio service mesh, and OPA Gatekeeper policy enforcement—supporting 200+ microservices across 15 development teams with zero-touch provisioning.
- **Full-Stack Observability:** Established full-stack observability platform (Prometheus, Grafana, Loki, Tempo, OpenTelemetry) with AI-powered anomaly detection, cutting P1 incidents by 42% and enabling distributed tracing across 30+ services.
- **Voice & Messaging Platform:** Delivered WhatsApp + Telegram AI Assistant platform using FastAPI, LangChain agents, and Redis caching—serving 10,000+ daily interactions with sub-500ms response times and 99.5% message delivery rate.
- **Compute Consolidation:** Administered VMware vSphere 7.x/8.x environment (ESXi, vCenter, vSAN) across Dell MX7000 and Cisco UCS blade infrastructure—managing 200+ VMs, achieving 40% compute consolidation ratio, and saving \$300K annually.
- **Patch & Change Governance:** Led Patch Management and Change Management for 500+ Windows Server instances using WSUS/SCCM and ServiceNow CAB process—maintaining 99.2% patch compliance and zero unplanned outages from change-related incidents.
- **Storage Architecture:** Designed NetApp SAN/NAS storage architecture with SnapVault/SnapMirror replication—providing 50TB+ enterprise storage with 99.99% availability and RPO/RTO of 15 min/1 hour for critical workloads.

Infrastructure Engineer · Trigent Software · Bangalore, India

Jun 2019 – Mar 2021

- **Virtualization Automation:** Managed Hyper-V and VMware virtualization clusters and Windows Server 2016/2019 environments for 300+ server estate—automating lifecycle management with PowerShell DSC scripts, reducing manual effort by 50%.
- **Enterprise Monitoring:** Implemented LogicMonitor enterprise monitoring across 400+ devices with custom alerting dashboards and escalation policies—achieving 95% alert signal-to-noise ratio improvement through AI-assisted threshold tuning.
- **Incident & Problem Management:** Executed end-to-end Incident and Problem Management via ServiceNow ITSM, maintaining SLA adherence at 97.5% for P1/P2 incidents and driving 30% reduction in recurring incident volume through root cause elimination.
- **Configuration as Code:** Automated server provisioning and configuration management using Ansible playbooks across 200+ Linux and Windows servers—reducing provisioning time from 4 hours to 20 minutes per server.
- **Azure Cloud Migration:** Led migration of 80+ on-premise workloads to Azure using Azure Migrate and Azure Site Recovery (ASR)—delivering cloud migration project 2 weeks ahead of schedule with zero production data loss.

KEY ARCHITECTURAL PROJECTS

Enterprise Agentic Operating System

LangGraph • FastAPI • Kubernetes • Redis • PostgreSQL • ArgoCD • OpenTelemetry

Multi-agent orchestration platform coordinating specialized AI agents (Planner, Executor, Critic, Memory) using LangGraph stateful cyclic graphs to handle 500+ concurrent enterprise workflows. Implemented persistent Redis and vector state with <120ms routing overhead.

Enterprise RAG Platform

LangChain • Weaviate • OpenSearch • FastAPI • Terraform • Grafana

Production RAG pipeline with hybrid semantic + BM25 search over 5M+ enterprise documents, metadata filtering, and re-ranking (45% retrieval precision boost). Integrated GraphRAG with Neo4j knowledge graph for entity traversal, reducing hallucinations by 3x.

AI Infrastructure Monitoring Copilot

Claude/GPT-4 • LogicMonitor API • ServiceNow • FastAPI • LangChain • MCP

SRE copilot with function-calling agents that ingest LogicMonitor telemetry, correlate alerts across topologies, and auto-generate enriched ServiceNow incidents with root cause hypotheses—cutting MTTR by 58% across 500+ servers.

Multi-Cloud Landing Zone & IaC Platform

Terraform • Pulumi • Ansible • GitHub Actions • Azure • AWS • GCP

Reusable Terraform module library (50+ modules) for multi-cloud Landing Zone provisioning. Automated GitOps PR validation with Checkov security scanning and Atlantis, establishing 100% compliance auditability.

WhatsApp & Voice AI Assistant Platform

FastAPI • LangChain • Whisper STT • ElevenLabs TTS • WebRTC • Redis

Multi-channel AI assistant supporting WhatsApp, Telegram, and WebRTC real-time voice calls. Processes 10,000+ daily interactions and 500+ voice receptionist calls/day with <1s end-to-end response latency.

CERTIFICATIONS & PROFESSIONAL DEVELOPMENT

COMPLETED	Microsoft Certified: Azure Administrator Associate (AZ-104)	IN PROGRESS	Microsoft Certified: Azure AI Engineer Associate (AI-102)
PLANNED	AWS Certified Solutions Architect – Associate	PLANNED	Google Professional Cloud DevOps Engineer

EDUCATION

Bachelor of Technology (B.Tech) — Electronics & Communication Engineering

Jawaharlal Nehru Technological University (JNTU), Anantapur

2013 – 2017